Hash Tables 1

Hash Tables

Objectives

- Understand the basic structure of a hash table and its associated hash function
  - Understand what makes a good (and a bad) hash function
- Understand how to deal with collisions
  - Open addressing
  - Separate chaining
- Be able to implement a hash table
- Understand how occupancy affects the efficiency of hash tables

Introduction

Problem Examples

- What can we do if we want rapid access to individual data items?
  - Looking up data for a flight in an air traffic control system
  - Looking up the address of someone making a 911 call
  - Checking the spelling of words by looking up each one in a dictionary
- In each case speed is very important
  - But the data does not need to be maintained in order

Possible Solutions

- Balanced binary search tree
  - Lookup and insertion in $O(\log n)$ time spoilers!
    - Which is relatively fast
  - Binary search trees also maintain data in order, which may be not necessary for some problems
- Arrays
  - Allow insertion in constant time, but lookup requires linear time
  - But, if we know the index of a data item lookup can be performed in constant time

Thinking About Arrays

- Can we use an array to insert and retrieve data in constant time?
  - **Yes** – as long as we know an item's index
- Consider this (very) constrained problem domain:
  - A phone company wants to store data about its customers in *Convenientville*
  - The company has approximately 9,000 customers
  - *Convenientville* has a single area code (555)

Living in Convenientville

- Create an array of size 10,000
  - Assign customers to array elements using their (four digit) phone number as the index
  - Only around 1,000 array elements are wasted
  - Customer data can be found in constant time using their phone numbers
- Of course this is not a general solution
  - It relies on having conveniently numbered *key* values

A (Poor) General Strategy

- In the Convientville example each possible key value was assigned an array element
  - With the index being the 4 digit phone number
  - Therefore the array size is the number of *possible* values
    10,000 in the example
    - Not the number of *actual* values
      9,000 in the example
- Consider two more examples that use this same general idea
  - Canadian phone numbers
  - Names

Phone Numbers in General

- Let's consider storing information about Canadians given their phone numbers
  - Between 000-000-0000 and 999-999-9999
- It's easy to convert phone numbers to integers
  - Just get rid of the "-"s
  - The keys range between 0 and 9,999,999,999
- Use *Convenientville* scheme to store data
  - But will this work?

A Really Big Array!

- If we use Canadian phone numbers as the index to an array how big is the array?
  - 9,999,999,999 (ten billion)
  - That's a really big array!
- An estimate of the current population of Canada is 35,623,680 source: [CIA World Fact Book]
  - That means that we will use around 0.3% of the array
    - That's a lot of wasted space
    - And the array may not fit in main memory …

Names

- What if we had to store data by name?
  - We would need to convert strings to integer indexes
- Here is one way to encode strings as integers
  - Assign a value between 1 and 26 to each letter
  - a = 1, z = 26 (regardless of case)
  - Sum the letter values in the string
- Not a very good method …

"dog" = 4 + 15 + 7 = 26

"god" = 7 + 15 + 4 = 26

Finding Unique String Values

- Ideally we would like to have a unique integer for each possible string
  - The "sum the letters" encoding scheme does not achieve this
- There is a simple method to achieve this goal
  - As before, assign each letter a value between 1 and 26
  - Multiply the letter's value by 26^i, where *i* is the position of the letter in the word:
    - "dog" = $4*26^2 + 15*26^1 + 7*26^0 = 3{,}101$
    - "god" = $7*26^2 + 15*26^1 + 4*26^0 = 5{,}126$

Afhahgm Vsyu

- The proposed system generates a unique integer for each string
  - But most strings are not meaningful
  - Given a string containing ten letters there are 26^{10} possible combinations of letters
    - Which gives 141,167,095,653,376 different possible strings
    - There are around 200,000 words in the English language

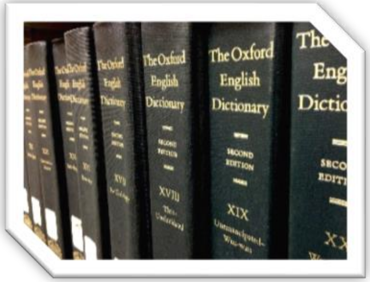

- It is not practical to create an array large enough to store all possible strings
  - Just like the general telephone number problem

So What's The Problem?

- In an ideal world we would know which key values were going to be recorded
  - The *Convenientville* example was close to ideal
- Most of the time this is not the case
  - Usually, key values are not known in advance
  - And, in many cases, the universe of possible key values is very large (e.g. names)
  - So it is not practical to reserve space for all possible key values

A Different Approach

- Don't determine the array size by the maximum possible number of keys
- Fix the array size based on the amount of data to be stored
  - Map the key value (phone number or name or some other data) to an array element
  - We will need to convert the key value to an integer index using a *hash function*
- This is the basic idea behind hash tables

Hash Tables

Hash Tables

- A *hash table* consists of an array to store data
  - Data often consists of complex types
    - Or pointers to such objects
  - An attribute of the object is designated as the table's *key*

- A *hash function* maps the key to an index
  - The key must first be converted to an integer
  - And mapped to an array index using a function
    - Often the modulo function

Collisions

- A hash function may map two different keys to the same index why?
  - Referred to as a *collision*
  - Consider mapping phone numbers to an array of size 1,000 where h = phone mod 1,000 this is not a good hash function …
    - Both 604-555-1987 and 512-555-7987 map to the same index (6,045,551,987 mod 1,000 = 987)
- A good hash function can significantly reduce the number of collisions
- It is still necessary to have a policy to deal with any collisions that may occur

Hash Functions

Hash Functions

- A hash function is a function that maps key values to array indexes
- Hash functions are performed in two steps
  - Map the key value to an integer
  - Map the integer to a legal array index
- Hash functions should have the following properties
  - Fast
  - Deterministic
  - Uniformity

Hash Function Speed

- Hash functions should be fast and easy to calculate
  - Access to a hash table should be nearly instantaneous and in constant time
  - Most common hash functions require a single division on the representation of the key
  - Converting the key to a number should also be able to be performed quickly

Deterministic Hash Functions

- A hash function must be *deterministic*
  - For a given input it must always return the same value
    - Otherwise it will not generate the same array index
    - And the item will not be found in the hash table
- Hash functions should therefore not be determined by
  - System time
  - Memory location
  - Pseudo-random numbers

Scattering Data

- A typical hash function usually results in some *collisions*
  - Where two different search keys map to the same index
  - A *perfect* hash function avoids collisions entirely
    - Each search key value maps to a different index
- The goal is to *reduce* the number and effect of collisions
- To achieve this the data should be distributed evenly over the table

Possible Values

- Any set of values stored in a hash table is an instance of the universe of possible values
- The universe of possible values may be much larger than the instance we wish to store
  - There are many possible combinations of 10 letters 26^{10}
  - But we might want a hash table to store just 1,000 names

Uniformity

- A good hash function generates each value in the output range with the same probability
  - That is, each legal hash table index has the same chance of being generated
- This property should hold for the universe of possible values *and for the expected inputs*
  - The expected inputs should also be scattered evenly over the hash table

A Bad Hash Function

- A hash table is to store 1,000 numeric estimates that can range from 1 to 1,000,000
  - Hash function is *estimate* % *n*
    - Where *n* = array size = 1,000
- Is the distribution of values from the universe of all possible values uniform?
- And what about the distribution of expected values?

Another Bad Hash Function

- A hash table is to store 676 names
  - The hash function considers just the first two letters of a name
    - Each letter is given a value where a = 1, b = 2, …
    - Function = (1st letter * 26 + value of 2nd letter) % 676
- Is the distribution of values from the universe of all possible values uniform?
- And what about the distribution of expected values?

General Principles

- Use the entire search key in the hash function
- If the hash function uses modulo arithmetic make the table size a prime number
- A simple and effective hash function is
  - Convert the key value to an integer, x
  - $h(x) = x \bmod tableSize$
    - Where *tableSize* is the first prime number larger than twice the size of the number of expected values

Caveat

- Consider mapping n values from a universe of possible values U into a hash table of size m
  - If $U \geq n \times m$
  - Then for any hash function there is a set of values of size n where all the keys map to the same location!
- Determining a good hash function is a complex subject
  - That is only introduced in this course

Converting Strings to Integers

Converting Strings to Integers

- A simple method of converting a string to an integer is to:
  - Assign the values 1 to 26 to each letter
  - Concatenate the binary values for each letter
    - Similar to the method previously discussed
- Using the string *cat* as an example:
  - c = 3 = 00011, a = 00001, t = 20 = 10100
  - So *cat* = 000110000110100 (or 3,124)
  - Note that $32^2 * 3 + 32^1 * 1 + 20 = 3{,}124$

Strings to Integers

- If each letter of a string is represented as a 32 bit number then for a length *n* string
  - value = $ch_0 * 32^{n-1} + \ldots + ch_{n-2} * 32^1 + ch_{n-1} * 32^0$
  - For large strings, this value will be very large
    - And may result in overflow
- This expression can be *factored*
  - $(\ldots(ch_0 * 32 + ch_1) * 32 + ch_2) * \ldots) * 32 + ch_{n-1}$
  - This technique is called *Horner's Method*
  - This minimizes the number of arithmetic operations
- Overflow can then be prevented by applying the *mod* operator after each expression in parentheses

Horner's Method Example

- Consider the integer representation of some string
  - $6*32^3 + 18*32^2 + 15*32^1 + 8*32^0$
  - $= 196{,}608 + 18{,}432 + 480 + 8 = 215{,}528$
- Factoring this expression results in
  - $(((6*32 + 18) * 32 + 15) * 32 + 8) = 215{,}528$
- Assume that this key is to be hashed to an index using the hash function *key* % 19
  - $215{,}528 \% 19 = 11$
  - $(((6*32 + 18) \% 19 * 32 + 15) \% 19 * 32 + 8) \% 19 = 11$
    - $210 \% 19 = 1$, and $47 \% 19 = 9$, and $296 \% 19 = 11$