


Graphs 2

Shortest Path



Shortest Path Problem

- What is the least cost path from one vertex to another?
 - Referred to as the shortest path between vertices
 - For weighted graphs this is the path that has the smallest sum of its edge weights
- Dijkstra's algorithm finds the shortest path between one vertex and all other vertices
 - The algorithm is named after its discoverer, Edgser Dijkstra



Shortest path between B and G ?

B-D-E-F-G

not



B-G or B-A-E-F-G

Dijkstra's Algorithm Overview

- Finds the shortest path to all vertices from the start vertex
- Performs a modified BFS that accounts for edge weights
 - Selects the node with the least cost from the start node
 - In an unweighted graph this reduces to a BFS
- In a weighted graph we need to assess the edge weights
 - And want choose the path with the least total cost
- Use a priority queue with an entry for each vertex
 - These entries record the cost of the path from the start to the vertex
 - Priority is given to lowest cost paths
 - Costs in the priority queue must be updated if a lower cost path is found
 - Efficiency issue if the priority queue is represented by a heap

BFS uses a queue

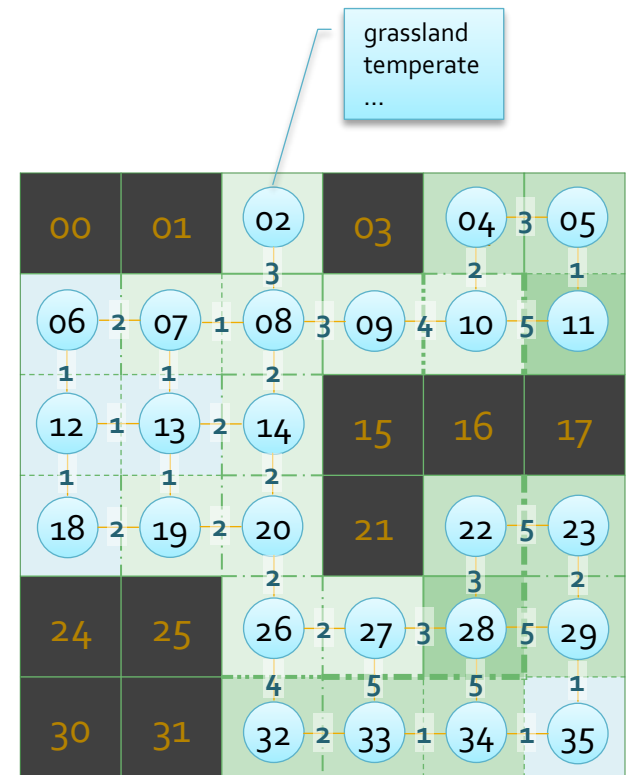
Example – Introduction

- Assume the grid represents a map
- Gray shaded areas are inaccessible
 - Mountains? 
 - Pits full of demons? 
- The cost to move from one location to the next varies between 1 and 5
 - Indicated by the thickness of the border between squares
- Presumably there is additional data about each area on the map
 - Occupants?
 - Terrain type?
 - Main export?

00	01	02	03	04	05
06	07	08	09	10	11
12	13	14	15	16	17
18	19	20	21	22	23
24	25	26	27	28	29
30	31	32	33	34	35

Example – Make a Graph

- Replace the map with a graph
 - Weighted and undirected
- Only need nodes for accessible regions
- The nodes contain only the information that is required for the application
- The edges record the edge weight
 - This graph is relatively sparse
 - So might choose to use an adjacency list
- We can now use the graph to answer questions about the application domain
- Such as – what is the shortest path between vertex 13 and all the other vertices?




Dijkstra's Algorithm – Initialization

- A record for each vertex is inserted into a priority queue, each record contains

- The vertex key (label)
- The cost to reach the vertex from the start
- The key of the previous vertex in the path

- These values are initially set as follows

- The cost to reach the start vertex is set to zero
- The cost to reach all other vertices is set to infinity and the parent vertex is set to the start
- Because the cost to reach the start vertex is zero it will be at the front of the priority queue



key	cost	parent
13	0	13
02	∞	13
04	∞	13
05	∞	13
...

Priority Queue Implementation

- Priority queues can be implemented with a heap
 - It is efficient for removing the highest priority item
 - In this case the element with the least cost
- Using a heap does have one drawback
 - Its elements will need to be accessed to update their costs
 - It is therefore useful to provide an index to its contents
- There are other data structures that can be used instead of a heap
 - That allow more efficient look-up than $O(n)$

Dijkstra's Algorithm – Main Loop


- Until the priority queue is empty
 - Remove the vertex with the least cost and insert it in a *results* list, making it the current vertex
 - The results list should be indexed by the search key of the vertices
 - Search in the adjacency list or matrix for vertices adjacent to the current vertex
 - For each such vertex, v
 - Compare the cost to reach v in the priority queue with the cost to reach v *via the current vertex*
 - If the cost via the current vertex is less – change v 's entry in the priority queue to reflect this new path

The results list contains {key, cost, parent} data and will need to be searched by key

Final Stage – Finding a Path


- When the priority queue is empty the results list contains all of the shortest paths from the start
- To find a path to a vertex look it up in the results list
 - The vertex's parent represents the previous vertex in the path
 - A complete path can be found by backtracking through all the parent vertices to the start vertex
 - A vertex's cost in the results list represents the total cost of the shortest path from the start to that vertex
 - Since the vertices need to be looked up by their key a hash table is a reasonable structure for the results list

Dijkstra's Algorithm Example 1



- Initialization
- Create priority queue
 - Entries for each vertex
 - label
 - cost to reach from start
 - previous vertex
- Create empty results list
 - Entries for vertices as priority queue
- Initialize costs in priority queue
 - 0 for start vertex
 - infinity for all others


Dijkstra's Algorithm Example 2



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

[illegible]


Dijkstra's Algorithm Example 3



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

[illegible]


Dijkstra's Algorithm Example 4



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

label	cost	parent	label	cost	parent
13	0	-			
07	1	13			
12	1	13			
19	1	13			
06	2	12			
08	2	07			
18	2	12			
14	2	13			


Dijkstra's Algorithm Example 5



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

label	cost	parent	label	cost	parent
13	0	-			
07	1	13			
12	1	13			
19	1	13			
06	2	12			
08	2	07			
18	2	12			
14	2	13			
20	3	19			
02	5	08			
09	5	08			
26	5	20			


Dijkstra's Algorithm Example 6



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

label	cost	parent	label	cost	parent
13	0	-	10	9	09
07	1	13	32	9	26
12	1	13	28	10	27
19	1	13	04	11	10
06	2	12	33	11	32
08	2	07			
18	2	12			
14	2	13			
20	3	19			
02	5	08			
09	5	08			
26	5	20			
27	7	26			

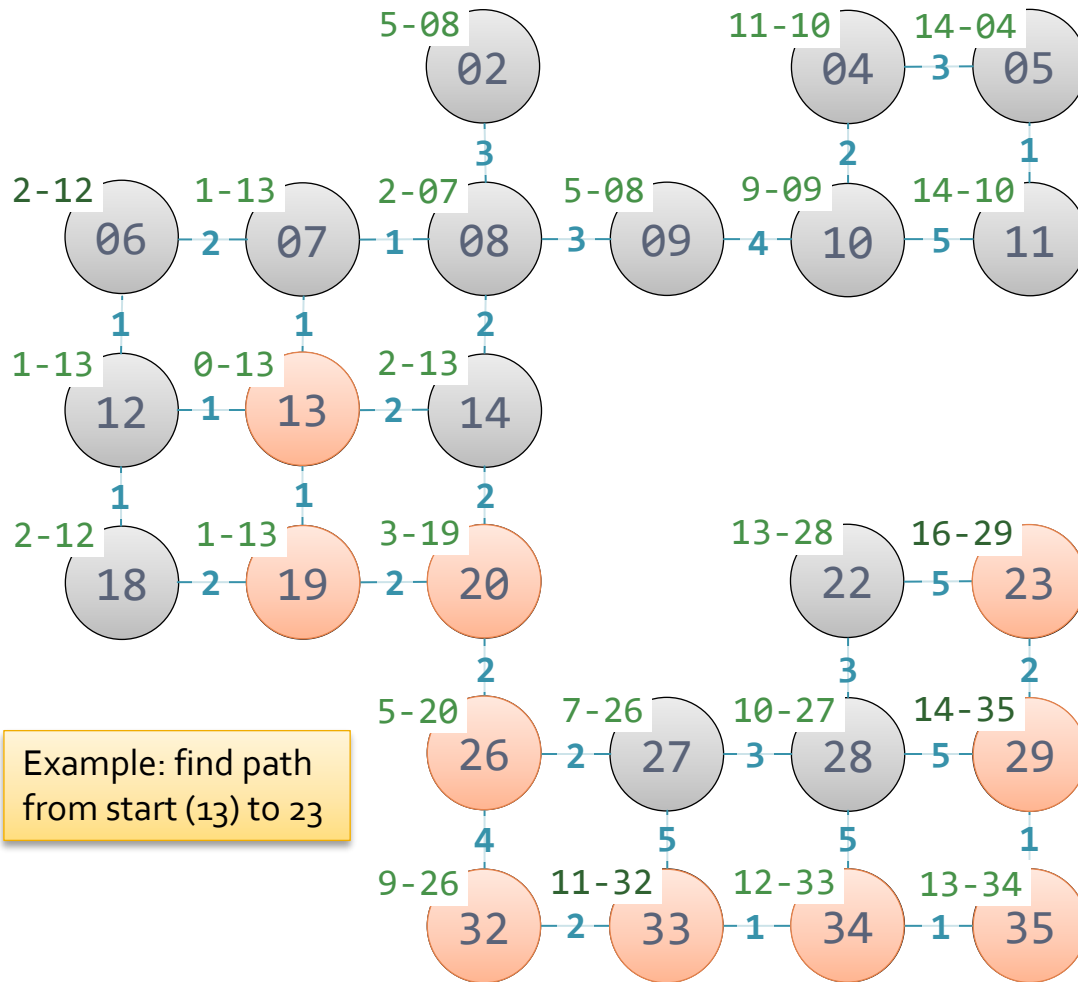
Dijkstra's Algorithm Example 7



- Remove vertex, v , from priority queue
 - Insert in results list
- Compare cost, c , to reach each adjacent vertex, u
 - If $c_u > (c_v + \text{edge}[v, u])$
 - Update cost and set u 's parent to v

label	cost	parent	label	cost	parent
13	0	-	10	9	09
07	1	13	32	9	26
12	1	13	28	10	27
19	1	13	04	11	10
06	2	12	33	11	32
08	2	07	34	12	33
18	2	12	22	13	28
14	2	13	35	13	34
20	3	19	05	14	04
02	5	08	11	14	10
09	5	08	29	14	35
26	5	20	23	16	29
27	7	26			

Retrieving the Shortest Path



- To find a path backtrack through results list
 - Start with end vertex
 - Find parent vertex
 - Repeat until start reached
 - Results list should allow efficient vertex searching

label	cost	parent	label	cost	parent
13	0	-	10	9	09
07	1	13	32	9	26
12	1	13	28	10	27
19	1	13	04	11	10
06	2	12	33	11	32
08	2	07	34	12	33
18	2	12	22	13	28
14	2	13	35	13	34
20	3	19	05	14	04
02	5	08	11	14	10
09	5	08	29	14	35
26	5	20	23	16	29
27	7	26			

Dijkstra's Algorithm Operations

- The cost of the algorithm depends on $|E|$ and $|V|$ and the data structures used in the algorithm
 - For the priority queue and results list
- Consider the process
 - Whenever a vertex is removed from the priority queue each of its adjacent edges must be found
 - There are $|V|$ vertices to be removed and
 - For each of $|E|$ edges it is necessary to
 - Retrieve the edge weight from the matrix or list
 - Look up the cost currently recorded in the priority queue for the edge's destination vertex

Dijkstra's Algorithm Analysis

- Assume a *heap* is used to implement the priority queue
- Building the heap takes $O(|V|)$ time
- Removing each vertex takes $O(\log|V|)$ time
 - For a total of $O(|V|*\log|V|)$
- Each of $|E|$ edges is processed once
 - Looking up and changing the current cost of a vertex in a heap takes $O(|V|)$ for an unindexed heap ($O(1)$ if the heap is indexed)
 - The heap property needs to be preserved after each change in cost for an additional cost of $O(\log|V|)$
 - Worst case cost is $|V| + |V|*\log|V| + |E|*(|V| + \log|V|)$
 - $O((|V|*\log|V|) + (|E|*|V|))$
 - If the heap is indexed the cost is $O((|V| + |E|) * \log|V|)$

The implementation we covered


Remember heaps do not support search

Bad!

Best, average case difference

What's an Indexed Heap?

- The cost analysis refers to an *indexed* heap
 - An index is a secondary data structure used to improve access to a primary data structure
 - Like an index ...
- Use a second data structure to find a heap node given the vertex label
 - The index records vertex labels and their corresponding indices in the underlying array of the heap
 - The index structure supports fast access by vertex label (its key)
 - Could use a hash table for the index
 - When an element is bubbled up in the heap the index also needs to be modified



Pathfinding with A*

- There are two drawbacks with Dijkstra's algorithm as a method of pathfinding
 - It finds paths from the start vertex to *all* other vertices, which results in wasted effort if only one path is required ... a trivial change ...
 - It only measures the cost *so far* - it does not *look ahead* to judge whether a path is likely to be a good one
- The A* algorithm addresses both these issues
 - It returns the path from the start vertex to the target vertex and
 - Uses an estimate of the remaining cost to reach the target to direct its search

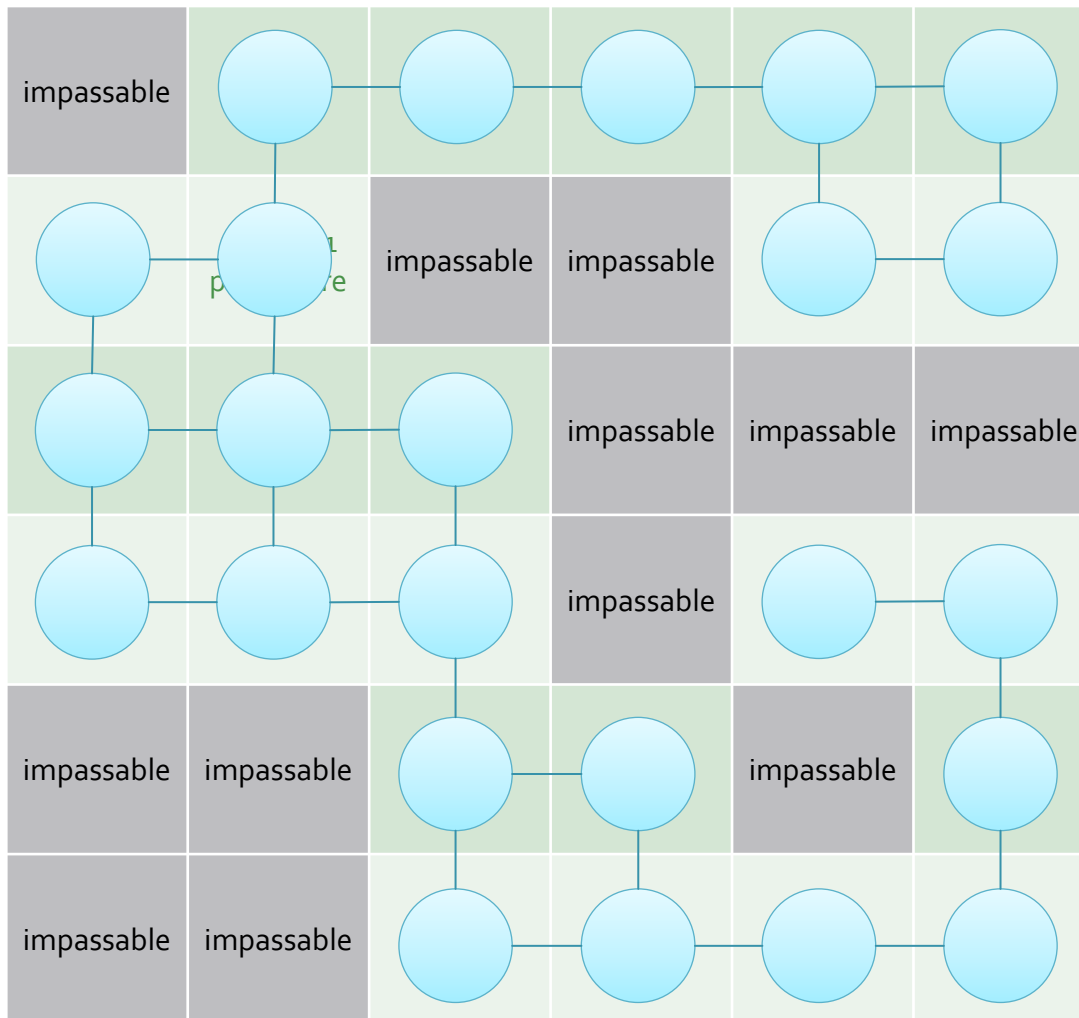
A* Algorithm

- The A* algorithm is similar to Dijkstra's algorithm
 - It performs a modified breadth first search and
 - Uses a priority queue to select vertices
- The A* algorithm uses a different cost metric, f , which is made up of two components
 - g – the cost to reach the current vertex from the start vertex (the same as Dijkstra's algorithm)
 - h – an estimate of the cost to reach the goal vertex from the current vertex
 - $f = g + h$

A* Heuristic – h

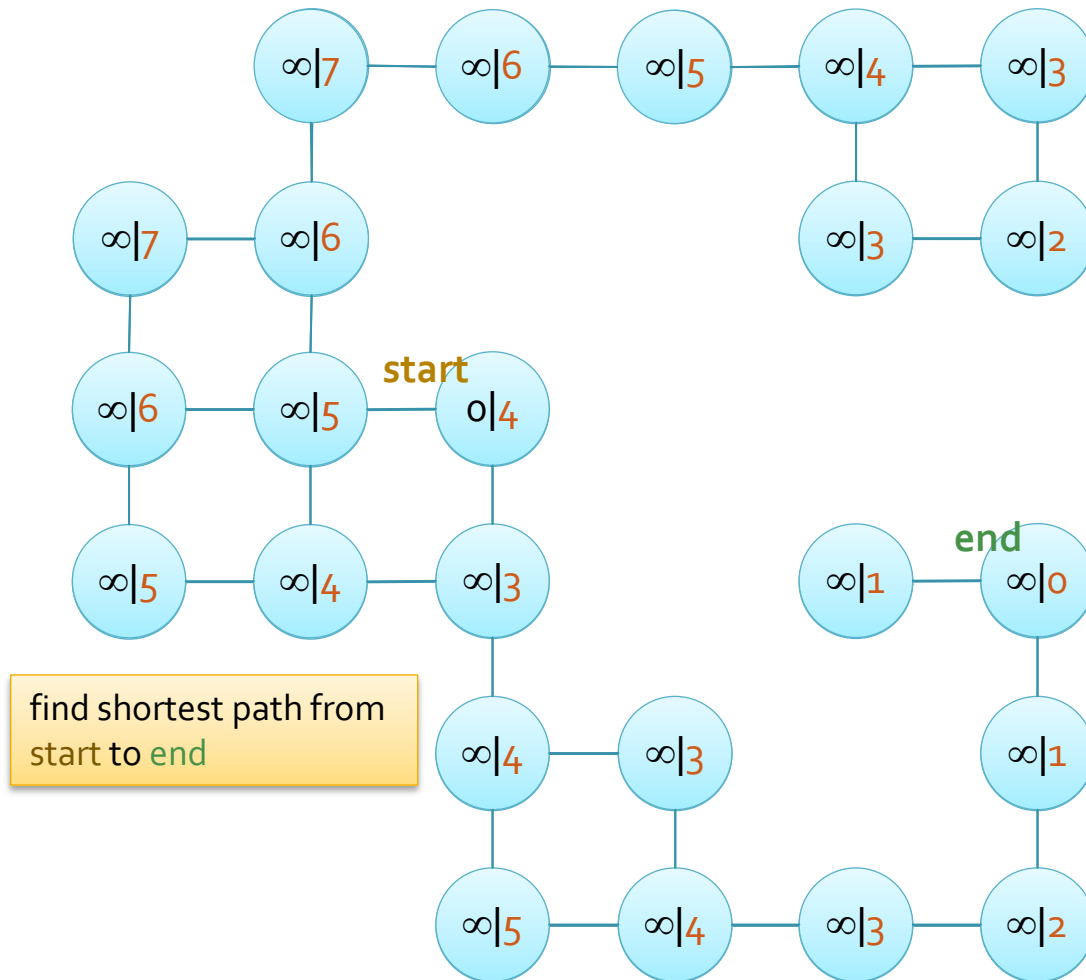
- The key to the efficiency of the A* algorithm is the accuracy of the heuristic, h
- To find an optimal path h should be *admissible*
 - It should not *overestimate* the cost of the path to the goal
 - Inadmissible heuristics may result in non-optimal paths
 - But may be faster than an inaccurate admissible heuristic
 - For a “good enough” solution it may be useful to use an inadmissible heuristic to speed up pathfinding
- If the heuristic is *perfect* the A* algorithm will find an optimal path with no backtracking

A* Search Example 1



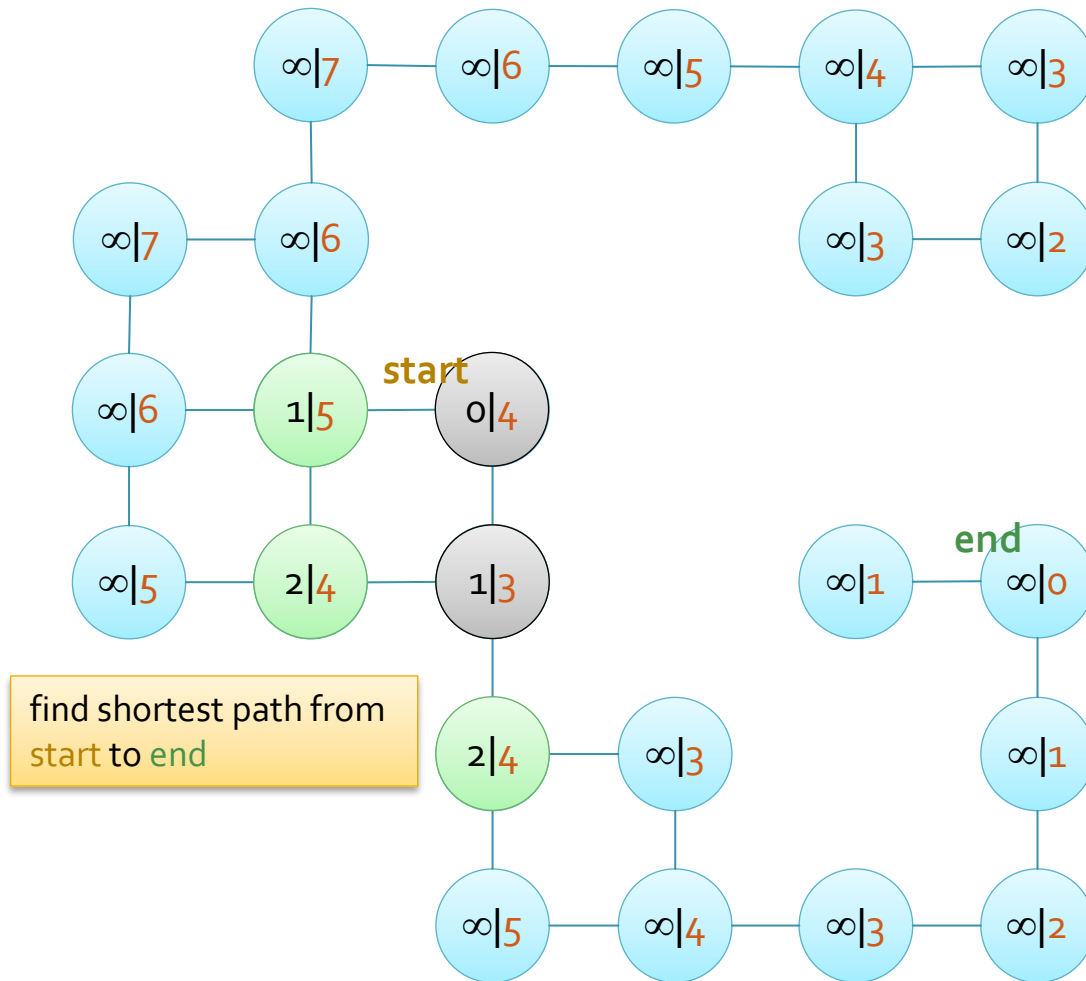
- Imagine a grid
 - gray squares
 - Gaps correspond to some game feature
- The corresponding graph is unweighted
 - Mostly to simplify the example
 - And focus on distinction between g and h
- g and h values
 - g – path length from start
 - h – number of squares to end square on full grid
 - straight-line cost
- Vertices annotated with g and h values
 - $g|h$
 - Grid labels not shown

A* Search Example 2



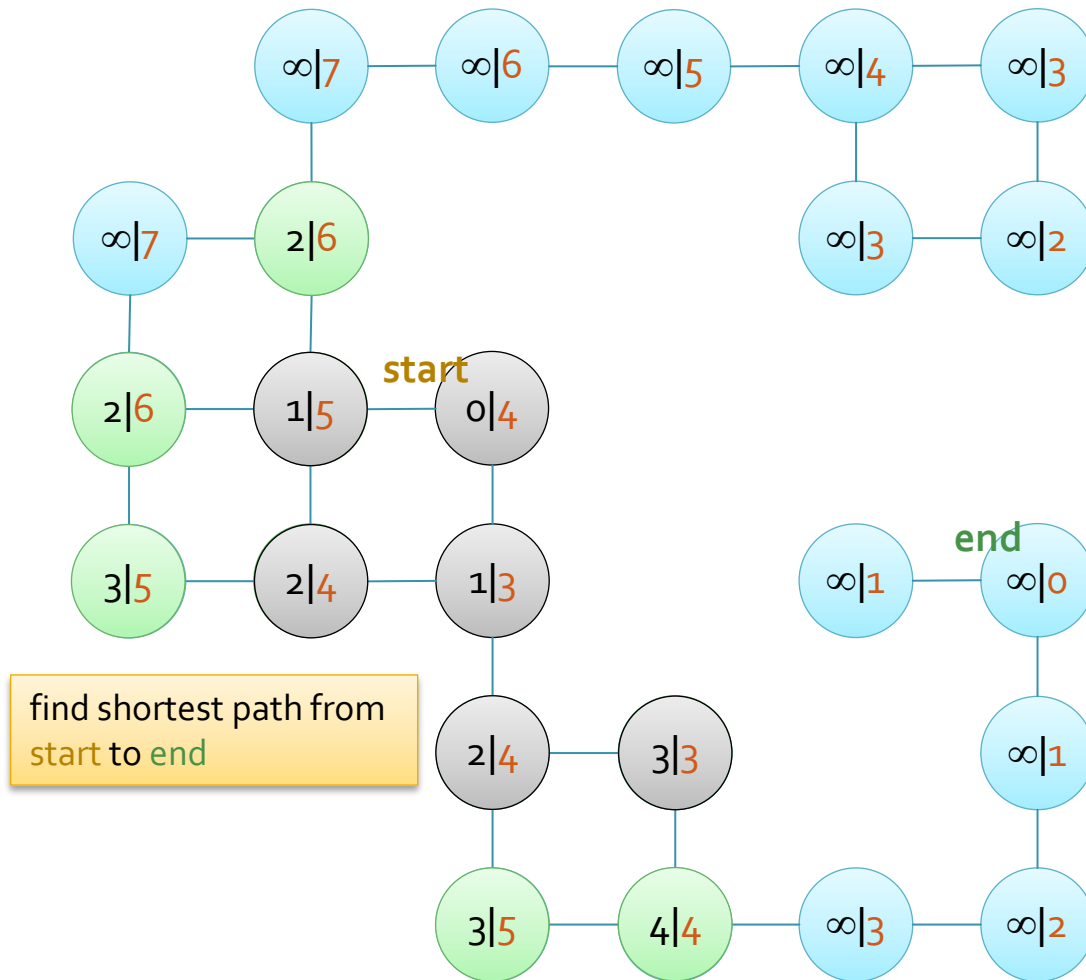
- The A* algorithm is similar to Dijkstra's except
 - Terminates when end vertex found
 - Values in priority queue based on f value
 - Which includes cost estimate to end (h)
- Initialization
 - Compute h for all vertices
 - Set g values
 - 0 for start
 - Infinity for all other vertices

A* Search Example 3



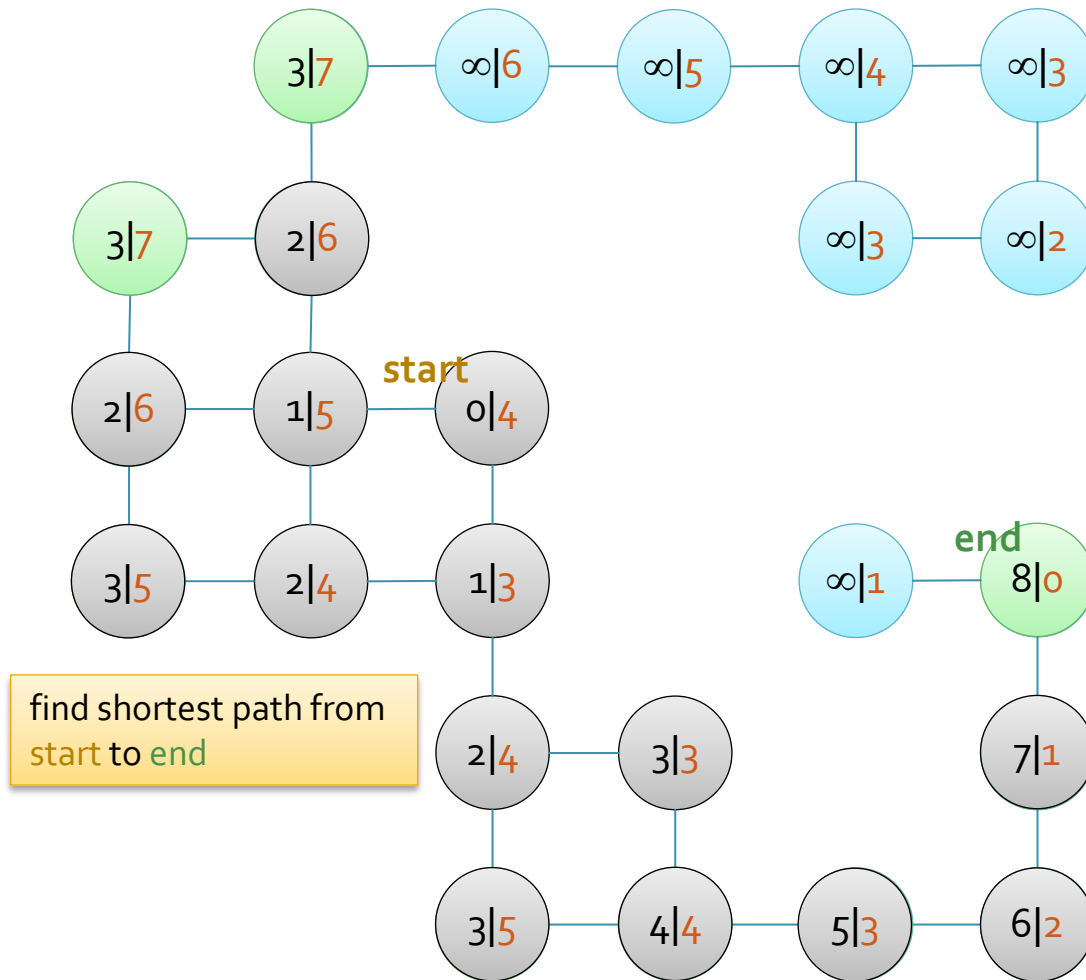
- Remove root from priority queue
 - Adjust g values of neighbours
- Remove new root
 - Remember – based on f , the sum of $g + h$
 - the vertex with $f = 4$

A* Search Example 4



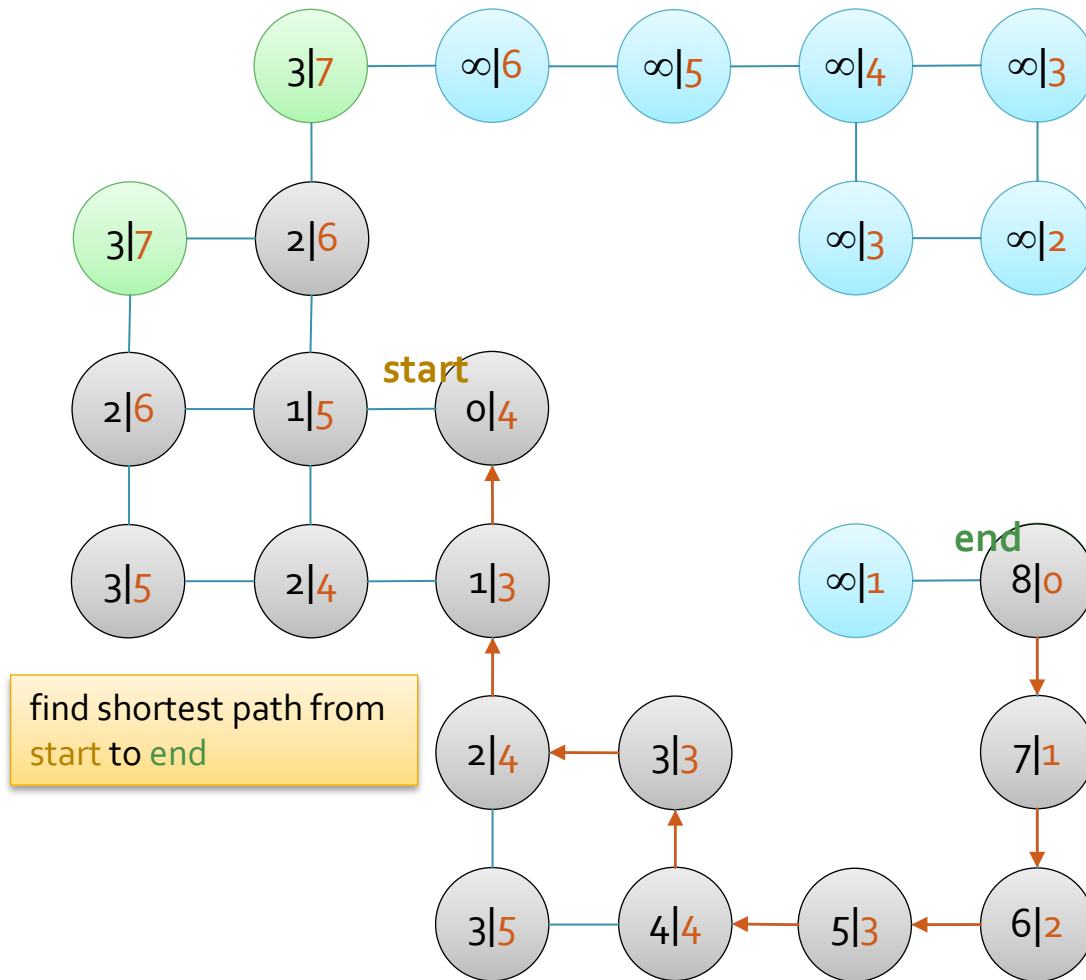
- Repeat
 - Until end removed from priority queue
- Note multiple vertices with $f = 6$
 - Remove whichever happens to be root

A* Search Example 5



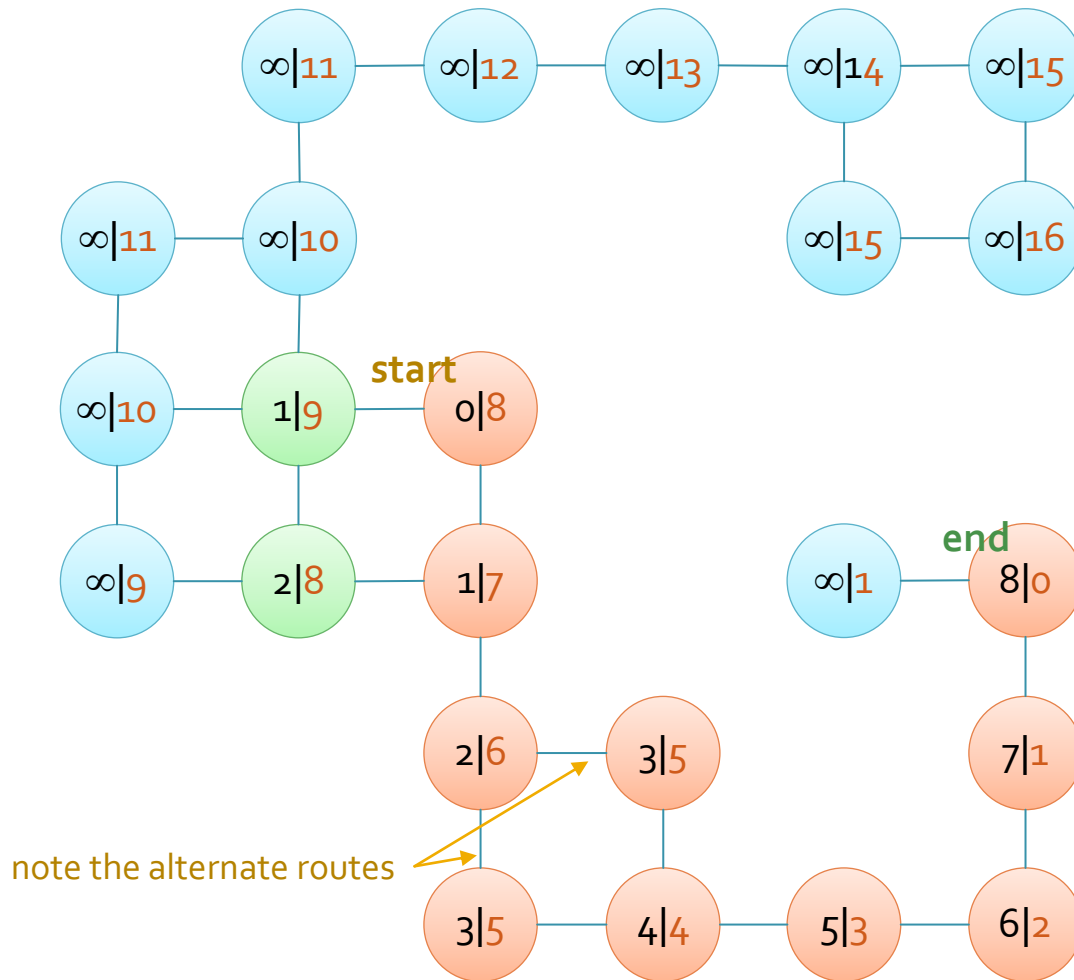
- Repeat
 - Until end removed from priority queue
- Now note multiple vertices with $f = 8$
 - Remove whichever happens to be root

A* Search Example 6



- Repeat
 - Until end removed from priority queue
- To find shortest path
 - Backtrack through results list
 - Like Dijkstra's algorithm

Perfect Heuristic



- A *perfect* heuristic gives correct cost
 - From vertex to end
 - Only vertices on shortest path are removed from the priority queue
- This graph shows a perfect heuristic
- But ...
 - There isn't a good general purpose heuristic generator
 - It helps to know about the graph
 - Bottlenecks
 - And other features